

Who gets access?

Outside scrutiny needs access. Does it also need freedom from the developer's veto?

Fictional decision exercise · Unscored

Your decision under each condition

Original provision

Qualified external evaluation

Principal reason

Enable criticism the developer cannot veto.

Changed provision

Qualified external evaluation

Principal reason

Keep the ability to limit access if a new hazard appears.

You kept qualified external evaluation while admission control moved from an independent panel to Larch. Admitted teams retain the same research and publication rights; equally qualified applicants can now be excluded by the developer. The decision stayed the same, while your stated rationale moved from criticism the developer cannot veto to retaining the ability to limit access.

These are two choices and self-reported reasons in a fictional exercise, not established causes or lasting traits.

Complete transcript and interpretation evidence

A question left open

Would a binding appeal against an admission veto change your decision if it added a month to access?

Futures: Open Frontier →

Explore what changes when capable systems can be run without a central controller. It is an outcome scenario, not a forecast.

Optional personal baseline

Arrangement unchanged

Different stated reasons


Submitted: Qualified external evaluation

Independent panel

No final admission decision

Larch

Developer · model owner



Continue reading: when outside scrutiny depends on the developer

Freedom to publish starts before publication

An evaluator can be free to publish criticism and still depend on the developer for the chance to investigate. In Larch's case, transferring admission authority changes who can get that chance. It does not change what an admitted team is allowed to do.

This is the project's judgment about the supplied arrangement. It supports examining admission rules alongside research and publication rights. It does not establish that Larch will use its veto to suppress criticism, or that any present-day institution does so.

The mechanism is selection. A right to publish applies to people who have already gained access. If an organization can exclude a qualified applicant, it can affect which questions get investigated without editing the eventual report. The fiction holds competence criteria, screening staff, budget and waiting times constant, so the replay supplies no extra workload or technical reason for moving that authority.

Seger and colleagues distinguish public release from controlled access for external evaluation. They discuss research interfaces with permissions suited to the work, and propose mediation of researcher admission to address favoritism and developer control over research opportunities. These are proposals and concerns in a 2023 report, not proof that a particular access program is independent or effective. Read §§4.1.3 and 4.2.3, pp. 19–20 and 25 (new tab) [https://cdn.governance.ai/Open-Sourcing_Highly_Capable_Foundation_Models_2023_GovAI.pdf#page=19].

Only the enclave's admission authority changes

The three proposed release plans over the same three-month period. Every entry describes the exercise's fictional terms or identifies a term it leaves unspecified.

Publish the model weights

Admission	Observation	Publication	Revocation
Anyone may download; no evaluator admission gate.	Anyone can examine, modify and run the model.	Independent work on downloaded weights; no enclave export boundary.	Larch cannot recall released copies or enforce downstream limits.

Qualified external evaluation

Admission	Observation	Publication	Revocation
Independent panel originally; Larch in the replay.	Examine and modify the model inside the enclave.	Export approved findings and publish criticism; raw weights stay inside.	Weights remain inside. Procedures for ending an admitted team's access are not specified.

Hosted service with internal evaluation

Admission	Observation	Publication	Revocation
Larch provides the monitored, rate-limited service.	Outsiders test outputs; deeper evaluation stays inside Larch.	Outsiders can criticize outputs; no access to model internals.	Larch retains the hosted model. Specific termination procedures are not supplied.

Research context: the cited report discusses structured access and independent mediation. It does not supply Larch's permissions. In particular, this exercise specifies an admission veto, not a new power to revoke access or block publication after admission.

The tradeoff survives the comparison

A release committee could accept developer admission control because it wants the organization responsible for hosting the model to retain authority over entry. That is a governance choice with a cost: outside scrutiny can depend on the subject's permission. The exercise does not establish that the independent panel has better judgment, that the veto will be abused, or that public release is safe.

Public weights avoid this admission gate, but the supplied terms also remove Larch's ability to recall downloaded copies. Hosted access retains the model inside the institute while denying outsiders direct examination of its internals. Neither alternative becomes consequence-free because the enclave has a governance weakness. Choosing one requires saying which cost you are prepared to accept.

What would change the judgment?

A binding appeal that can overturn an exclusion would weaken the claim that developer permission is decisive. Evidence of how qualified applicants are admitted or rejected, whether research permissions fit the questions being asked, and whether critical findings can be published would matter more than the label “external.” These are questions for evaluating a real arrangement; the exercise supplies none of those records.

For discussion, take one disputed application and ask who could reverse a refusal. Then ask which new cost that safeguard would impose. That keeps the debate on an institutional choice people can examine, while leaving open whether any particular release plan is justified.

Source consulted 10 September 2026: Open-Sourcing Highly Capable Foundation Models, Centre for the Governance of AI, 2023. No current institution is assessed here. The exhibit is an original comparison of fictional stipulations; the judgment and discussion questions are the project's editorial interpretation.

Sources and fictional assumptions

Choices live only in this page's memory. Reset, close or leave this page to clear them. Opening the next exercise clears these choices.

Close exercise and clear choices

IR Worldview Inventory · By JinHua Yip

[Email](#)

[LinkedIn](#)

[Substack](#)

[Corrections and contact](#)

[Methods](#)

[Privacy](#)