

Precautionary Steward

FRONTIER-RISK FIRST · DEPLOYMENT-FIRST · COORDINATION-FIRST

Oversight leans toward stronger public supervision rather than developer discretion. The legitimacy score does not resolve a disagreement between expertise and public authorization. These positions leave distinct questions about who can enforce a rule and who can authorize it. Neither score tells us how an appeal should work.

An interpretation of the recorded scores. This link does not contain an answer transcript; the example below is an application of the reasoning, not a forecast of your choice.

Who can stop a deployment?

Suppose a regulator and a developer disagree over a failed safety evaluation. A one-month delay would disrupt a useful public service, and the regulator's technical team is small.

Applied here, the supervision argument gives the regulator an enforceable stop power, with access to outside technical evidence. Its legitimacy still depends on the grounds for intervention and a way to challenge errors.

The strongest objection

A poorly equipped regulator can become dependent on the firms it oversees. Formal authority may then slow useful deployment without producing independent judgment.

What would change the argument. The argument changes if the regulator can obtain independent expertise, publish its reasons and resolve appeals before delay defeats the service's purpose.

An illustrative application of the reasoning. This is not a recorded answer or a prediction.

Positions furthest from the midpoint

Oversight leans toward stronger public supervision rather than developer discretion.

Risk horizon



Severe frontier hazards carry substantial weight in this reading.

Openness vs control



Controlled access carries more weight than broad diffusion.

The marker locates the recorded position on a 1–7 scale. Either endpoint can express a strong position; this is not a progress bar.

Nearest alternative: State Capacity Builder

DIMENSION	PRECAUTIONARY STEWARD	STATE CAPACITY BUILDER	YOU
Risk horizon Present harms — Frontier risk	Frontier risk	No strong model emphasis	5.6 Comparison term: 1.440
Deployment pace Ship and learn — Slow on evidence	Slow on evidence	No strong model emphasis	3.5 Comparison term: -0.450
Military role Keep AI out — Bounded defense use	Keep AI out	No strong model emphasis	4.6 Comparison term: -0.540

These are authored model emphases, not measured expectations about people.

The recorded positions favor Precautionary Steward over State Capacity Builder most through risk horizon.

A case that distinguishes these emphases. Would a poorly quantified catastrophic hazard justify a delay when deployment already reduces a documented present-day harm?

Each term is the recorded axis position minus 4, multiplied by the difference between these two frozen model weights. Positive terms favor Precautionary Steward; negative terms favor State Capacity Builder.

What would change this

Who should have the final say when an expert panel recommends release but an accountable public body refuses it?

A new question to consider. No answer is recorded or rescored here.

[Read the archetype page](#)

[Take the IR Foundation](#)

[View Profile](#)

IR RELATIONSHIP

AI and the Foundation remain separate domain reads

No IR Foundation baseline is saved on this device yet. The AI result still stands on its own.

Foundation baseline not yet saved

You may save a Foundation result and read the two records side by side inside Profile. The records remain separate unless a reviewed bridge is explicitly authored.

[Take the Foundation](#) →

Full analysis

Questions raised by these positions

Frontier release thresholds

What evidence should trigger release, delay, staged access, or a stronger stop?

The pace score does not establish a general release or delay rule. Severe frontier hazards carry substantial weight in this reading.

Oversight and enforcement

Who should have authority when frontier developers, agencies, and outside evaluators disagree?

Oversight leans toward stronger public supervision rather than developer discretion. The legitimacy score does not resolve a disagreement between expertise and public authorization.

International order

When should cooperation be pursued despite rivalry, and when does verification need to come first?

Coordination receives more weight than preserving competitive advantage on this axis.

Openness and diffusion

Where does broader access reduce dependence, and where does it outrun governance?

Controlled access carries more weight than broad diffusion.

What the remaining positions say

Military role

The military-use score does not identify which specific uses would be acceptable.

Human future

This mixed axis cannot separate views on transformation, human agency and particular imagined futures.

The Precautionary Steward reference argument

The burden of proof should rise with capability, especially where loss of control, catastrophic misuse, or systemic dependence could follow.

These are the authored category's proposals and challenges. Individual positions above can depart from them; the category does not supply missing answers.

Institutional priorities

- credible thresholds and eval gates
- stronger deployment restraint under uncertainty
- institutional willingness to pause when warning signs accumulate

Evidence that would challenge this argument

- Robust evidence that controlled real-world deployment reduces risk faster than withholding systems.
- Strong evidence that rivals would systematically exploit unilateral restraint without meaningful verification.
- Repeated cases where precautionary institutions fail to act competently enough to justify the delays they impose.


Read the reference argument, its objections and sources

Saved category wording

The registered model describes Precautionary Steward in the words below. This is a category reference; the recorded positions and any co-leading readings determine how far it applies.

You center severe frontier risks, prefer threshold-based caution, and are comparatively willing to accept slower deployment in exchange for stronger control.

Your positions on the AI questions











Each bar shows distance from the midpoint of the 1-7 axis. Bar length does not measure rarity or isolate how much one axis caused the archetype result. Methods sets out the limits →

Model tuple: bank 3, scorer 2. Saved comparison: Precautionary Steward minus State Capacity Builder = 0.50. Rounding residual: 0.0000.

- Risk horizon: 1.440
- Public oversight: -0.720
- Openness vs control: 0.560
- Military role: -0.540
- Competition vs coordination: 0.520
- Legitimacy and rule-setting: -0.490
- Deployment pace: -0.450
- Human future: 0.180

Axis profile



Read the result from another angle

Who should have the final say when an expert panel recommends release but an accountable public body refuses it? A poorly equipped regulator can become dependent on the firms it oversees. Formal authority may then slow useful deployment without producing independent judgment. These sources offer ways to examine that question; their authors do not endorse this result or the scenario.

CHALLENGE YOUR VIEW

Which assumptions in this reading deserve the hardest challenge?
These sources supply competing arguments, not predictions of your views.

On the Dangers of Stochastic Parrots

Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, Shmargaret Shmitchell · 2021
The canonical present-harms critique of scale-first language-model development.

Harms of AI

Daron Acemoglu · 2021
A strong challenge to governance approaches that assume more AI or more state support is obviously beneficial.

Atlas of AI

Kate Crawford · 2021
A material critique of AI as an extractive infrastructure, not just a software or safety problem.

Open-Sourcing Highly Capable Foundation Models

Elizabeth Seger et al. · 2023
Useful when testing openness-first instincts against the strongest misuse and proliferation critique.

EXPLORE THIS QUESTION

Cross-check in IR modules — Compare the AI result with your Security and Technology answers.

READ THE PROFILE’S FOUNDATIONS

Shared field maps and baseline texts before you settle into one governance camp.

AI Governance: A Research Agenda

Allan Dafoe · 2018
Still one of the clearest maps of the field: alignment, concentration, institutional design, misuse, and global governance.

International AI Safety Report 2025

Useful as a shared scientific baseline for advanced-AI safety debates across countries rather than a single camp's framing.

OECD AI Principles

OECD · 2019 / updated 2023

A high-level intergovernmental baseline for trustworthy AI, rights, robustness, and accountability.

NIST AI RMF: Generative AI Profile

NIST · 2024

Shows how governance looks when translated into operational risk-management language.

EXPLORE THIS QUESTION

Browse the AI atlas — See nearby archetypes and the reading shelves that sit beside them.

Open the AI field guide — Read the axes, scope limits, and under-modeled perspectives behind the result.

GO DEEPER

Operational frameworks, official documents, and live policy tools that show how these debates cash out in practice.

Frontier AI Safety Policies

METR · 2026

Useful for comparing how concrete threshold and mitigation regimes differ across frontier actors.

Managing AI Risks in an Era of Rapid Progress

Yoshua Bengio et al. · 2023

A broad coalition statement capturing the frontier-risk mainstream case for stronger governance.

Common Elements of Frontier AI Safety Policies

METR · 2025 / 2026 site updates

Synthesizes what labs are actually doing around thresholds, model evaluations, weight security, and deployment mitigations.

Responsible Scaling Policy

Anthropic · 2026 update

A living example of how a frontier lab publicly structures capability thresholds and safeguards.

Updated Preparedness Framework

OpenAI · 2025

Useful for comparing another frontier-lab approach to severe-risk measurement and mitigation.

Safety Cases at AISI

AI Security Institute · 2026

Pushes the debate from generic principles toward structured arguments that a system is safe in a given context.

EXPLORE THIS QUESTION

- Read methods — See the scope, guardrails, and limits behind the whole project.
- Open references — Use the broader bibliography when you want a wider shelf than the result page alone.

WIDEN THE FRAME

How these debates look beyond U.S. and European defaults, including Chinese and broader global-governance lenses.

State of AI Safety in China

Concordia AI · 2025

The best single English-language guide to the diversity and maturation of Chinese AI safety discourse.

Global AI Governance Initiative

Ministry of Foreign Affairs of the PRC · 2023

Essential for understanding China's official framing around sovereignty, development, safety, and multilateral governance.

Global AI Governance Action Plan

2025 World AI Conference and High-Level Meeting on Global AI Governance · 2025

A current official text on cooperation, development, and global-governance architecture beyond a U.S.-centric frame.

Recommendation on the Ethics of Artificial Intelligence

UNESCO · 2021

Useful for human-rights, dignity, and global ethics framing that extends beyond frontier-lab discourse.

EXPLORE THIS QUESTION

- Keep browsing the shelf — Move from these entry points into the full bibliography when you want more range.

Coverage limits

The Compass covers a defined set of AI governance debates and leaves others outside its scope. The raw scores locate this response inside the model. The archetype is the closest fit among six authored profiles. Full methods note → AI scope →

Report a factual problem →

RESEARCH COLLECTION

Research responses are not collected

Draft answers and saved AI Governance Compass history stay in this browser. Generated result links contain the data needed to reopen a result; Focus Area links include selected answer IDs. Separate first-party aggregate counters may receive derived buckets when coarse measurement is enabled. The site does not collect research responses or contact details.

Any future research study would require separate information about its purpose, consent, retention, deletion, and access rules.

[Read the privacy and data-use note.](#)