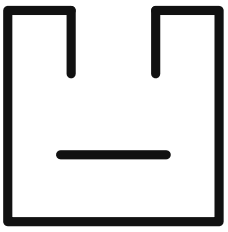
AI GOVERNANCE



Democratic Guardrailist

PRESENT-HARMS FIRST · DEPLOYMENT-FIRST · COMPETITIVE HEDGER

Controlled access carries more weight than broad diffusion. Present harms carry more weight than distant frontier hazards on this axis. The case for access controls here need not rest on catastrophe: present misuse, accountability or market power could also motivate limits. The scores cannot distinguish those reasons.

An interpretation of the recorded scores. This link does not contain an answer transcript; the example below is an application of the reasoning, not a forecast of your choice.

A frontier model with disputed misuse evidence

Suppose an independent team reports a serious misuse capability. Other researchers cannot yet reproduce the finding, and the developer controls admission to its testing service.

Applied here, the control-oriented argument is to retain restrictions while independent investigators test the reported hazard. Developer-only review would leave a conflict of interest even if the restriction itself were justified.

The strongest objection

The strongest objection is that the gatekeeper can suppress inconvenient findings or protect incumbents. Restriction can make the evidence needed to justify release harder to obtain.

What would change the argument. The argument changes if independent reproduction shows that the proposed access channel either contains the hazard or reliably enables it.

An illustrative application of the reasoning. This is not a recorded answer or a prediction.

Positions furthest from the midpoint

Openness vs control 5.8



Open diffusion — Staged access

Controlled access carries more weight than broad diffusion.

Public oversight 5.3



Lab-led — Public supervision

Oversight leans toward stronger public supervision rather than developer discretion.

Risk horizon 3.0



Present harms — Frontier risk

Present harms carry more weight than distant frontier hazards on this axis.

Each marker is your position on a 1 to 7 scale between two poles. Either end is a strong position.

Nearest alternative: State Capacity Builder

DIMENSION	DEMOCRATIC GUARDRAILIST	STATE CAPACITY BUILDER	YOU
Legitimacy and rule-setting Expert rule-setting — Public mandate	Public mandate	Expert rule-setting	4.8 Largest supporting term: 1.200

DIMENSION	DEMOCRATIC GUARDRAILIST	STATE CAPACITY BUILDER	YOU
Deployment pace Ship and learn — Slow on evidence	Slow on evidence	No strong model emphasis	3.1 Largest opposing term: -0.450

These are authored model emphases, not measured expectations about people.

The largest supporting term is Legitimacy and rule-setting. The largest opposing term is Deployment pace.

A case that distinguishes these emphases. Who should authorize a consequential release when technical experts and an accountable public body reach different conclusions?

▼ How the comparison terms are calculated

Each term is your position on the axis, measured from the midpoint, multiplied by how differently the two approaches weight that axis. Positive terms favor Democratic Guardrailist; negative terms favor State Capacity Builder.

What would change this

Would you allow independent research access if the developer could veto admission, while public release remained irreversible?

Nothing you decide here is recorded or rescored.

Read the archetype page | Take the IR Foundation | View Profile

Connect this with another question

Profile connects available choices about access, public capacity and cooperation. With one saved subject, it offers a question to examine next.

Open Profile · Compare results

Which futures would you want?

Browse nineteen possible worlds and say which conditions you would want in them. Futures is a separate exercise and starts empty.

Reasoning and policy debates

▼ Policy debates behind the approaches

Questions raised by these positions

Frontier release thresholds

What evidence should trigger release, delay, staged access, or a stronger stop?

The pace score does not establish a general release or delay rule. Present harms carry more weight than distant frontier hazards on this axis.

Oversight and enforcement

Who should have authority when frontier developers, agencies, and outside evaluators disagree?

Oversight leans toward stronger public supervision rather than developer discretion. The legitimacy score does not resolve a disagreement between expertise and public authorization.

International order

When should cooperation be pursued despite rivalry, and when does verification need to come first?

The geopolitical score leaves room for both coordination and strategic competition; it does not specify which comes first in a conflict.

Openness and diffusion

Where does broader access reduce dependence, and where does it outrun governance?

Controlled access carries more weight than broad diffusion.

What the remaining positions say

Military role

The military-use score does not identify which specific uses would be acceptable.

Human future

This mixed axis cannot separate views on transformation, human agency and particular imagined futures.

The Democratic Guardrailist reference argument

AI should not be governed only by labs, security agencies, or technical insiders; public legitimacy and rights protections have to remain visible.

These are the authored category's proposals and challenges. Individual positions above can depart from them; the category does not supply missing answers.

Institutional priorities

- public accountability and contestability
- stronger oversight for concentrated private power
- rights-conscious governance of deployment and monitoring

Evidence that would challenge this argument

- Strong evidence that public institutions cannot build or hire enough expertise to supervise credibly.
- Evidence that certain technocratic or lab-led mechanisms outperform public processes without becoming self-serving.
- Repeated cases where formal accountability structures exist on paper but fail to shape actual deployment behavior.

Read the reference argument, its objections and sources

▼ Saved category wording

The registered model describes Democratic Guardrailist in the words below. This is a category reference; the recorded positions and any co-leading readings determine how far it applies.

You want strong public rules, visible accountability, and governance answerable to democratic and civic institutions as well as frontier labs.

Your positions on the AI questions

TOWARD THE LOW POLE	MODEL MIDPOINT	TOWARD THE HIGH POLE	
Openness vs control		Staged access	5.8
Public oversight		Public supervision	5.3
Risk horizon		Present harms	3.0
Deployment pace		Ship and learn	3.1
Legitimacy and rule-setting		Public mandate	4.8
Human future		Human control	4.3
Military role		Bounded defense use	4.2
Competition vs coordination		Rivalry	4.1

Each bar shows distance from the midpoint of the 1-7 axis. Bar length does not measure rarity or isolate how much one axis caused the archetype result. Methods sets out the limits →

Model tuple: bank 3, scorer 2. Saved comparison: Democratic Guardrailist minus State Capacity Builder = 0.89. Rounding residual: 0.0000.

- Legitimacy and rule-setting: 1.200
- Deployment pace: -0.450
- Openness vs control: 0.180
- Military role: -0.140
- Public oversight: 0.130
- Human future: 0.120
- Risk horizon: -0.100
- Competition vs coordination: -0.050

Axis profile

Risk horizon 3



Present harms — Frontier risk

Present harms carry more weight than distant frontier hazards on this axis.

Deployment pace 3.1



Ship and learn — Slow on evidence

The pace score does not establish a general release or delay rule.

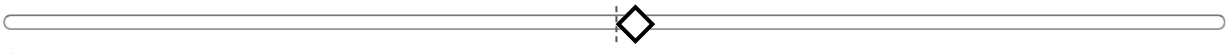
Public oversight 5.3



Lab-led — Public supervision

Oversight leans toward stronger public supervision rather than developer discretion.

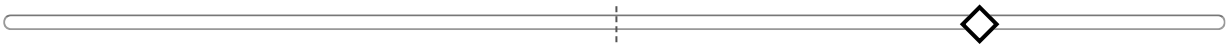
Competition vs coordination 4.1



Coordination — Rivalry

The geopolitical score leaves room for both coordination and strategic competition; it does not specify which comes first in a conflict.

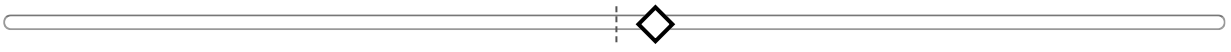
Openness vs control 5.8



Open diffusion — Staged access

Controlled access carries more weight than broad diffusion.

Military role 4.2



Keep AI out — Bounded defense use

The military-use score does not identify which specific uses would be acceptable.

Legitimacy and rule-setting 4.8



Expert rule-setting — Public mandate

The legitimacy score does not resolve a disagreement between expertise and public authorization.

Human future 4.3



Accept transformation — Human control

This mixed axis cannot separate views on transformation, human agency and particular imagined futures.

Read the result from another angle

Would you allow independent research access if the developer could veto admission, while public release remained irreversible? The strongest objection is that the gatekeeper can suppress inconvenient findings or protect incumbents. Restriction can make the evidence needed to justify release harder to obtain. These sources offer ways to examine that question; their authors do not endorse this result or the scenario.

CHALLENGE YOUR VIEW

Which assumptions in this reading deserve the hardest challenge? These sources supply competing arguments, not predictions of your views.

When Speed Kills: Lethal Autonomous Weapon Systems and the Dangers of Rushing to Weaponize AI

Michael Horowitz · 2019

A strong security-focused critique of governance that ignores the pace pressures created by military competition.

Incompleteness: A Regulatory Design Challenge for AI

Gillian Hadfield · 2019

A useful challenge to the idea that more democratic procedure automatically yields governance that can keep up with fast-moving systems.

On the Dangers of Stochastic Parrots

Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, Shmargaret Shmitchell · 2021

The canonical present-harms critique of scale-first language-model development.

Harms of AI

Daron Acemoglu · 2021

A strong challenge to governance approaches that assume more AI or more state support is obviously beneficial.

Atlas of AI

Kate Crawford · 2021

A material critique of AI as an extractive infrastructure, not just a software or safety problem.

Open-Sourcing Highly Capable Foundation Models

Elizabeth Seger et al. · 2023

Useful when testing openness-first instincts against the strongest misuse and proliferation critique.

EXPLORE THIS QUESTION

Cross-check in IR modules — Compare the AI result with your Security and Technology answers.

READ THE PROFILE'S FOUNDATIONS

Shared field maps and baseline texts before you settle into one governance camp.

AI Governance: A Research Agenda

Allan Dafoe · 2018

Still one of the clearest maps of the field: alignment, concentration, institutional design, misuse, and global governance.

International AI Safety Report 2025

Independent international expert group · 2025

Useful as a shared scientific baseline for advanced-AI safety debates across countries rather than a single camp's framing.

OECD AI Principles

OECD · 2019 / updated 2023

A high-level intergovernmental baseline for trustworthy AI, rights, robustness, and accountability.

NIST AI RMF: Generative AI Profile

NIST · 2024

Shows how governance looks when translated into operational risk-management language.

EXPLORE THIS QUESTION

Browse the AI atlas — See nearby archetypes and the reading shelves that sit beside them.

Open the AI field guide — Read the axes, scope limits, and under-modeled perspectives behind the result.

GO DEEPER

Operational frameworks, official documents, and live policy tools that show how these debates cash out in practice.

Common Elements of Frontier AI Safety Policies

METR · 2025 / 2026 site updates

Synthesizes what labs are actually doing around thresholds, model evaluations, weight security, and deployment mitigations.

Responsible Scaling Policy

Anthropic · 2026 update

A living example of how a frontier lab publicly structures capability thresholds and safeguards.

Updated Preparedness Framework

OpenAI · 2025

Useful for comparing another frontier-lab approach to severe-risk measurement and mitigation.

Safety Cases at AISI

AI Security Institute · 2026

Pushes the debate from generic principles toward structured arguments that a system is safe in a given context.

EXPLORE THIS QUESTION

Read methods — See the scope, guardrails, and limits behind the whole project.

Open references — Use the broader bibliography when you want a wider shelf than the result page alone.

WIDEN THE FRAME

How these debates look beyond U.S. and European defaults, including Chinese and broader global-governance lenses.

State of AI Safety in China

Concordia AI · 2025

The best single English-language guide to the diversity and maturation of Chinese AI safety discourse.

Global AI Governance Initiative

Ministry of Foreign Affairs of the PRC · 2023

Essential for understanding China's official framing around sovereignty, development, safety, and multilateral governance.

Global AI Governance Action Plan

2025 World AI Conference and High-Level Meeting on Global AI Governance · 2025

A current official text on cooperation, development, and global-governance architecture beyond a U.S.-centric frame.

Recommendation on the Ethics of Artificial Intelligence

UNESCO · 2021

Useful for human-rights, dignity, and global ethics framing that extends beyond frontier-lab discourse.

EXPLORE THIS QUESTION

Keep browsing the shelf — Move from these entry points into the full bibliography when you want more range.

Coverage limits

The Compass covers a defined set of AI governance debates and leaves others outside its scope. The raw scores locate this response inside the model. The archetype is the closest fit among six authored profiles. Full methods note → AI scope →

Report a factual problem →

AI approaches | AI Field Guide

RESEARCH COLLECTION

Research responses are not collected

Draft answers and saved AI Governance Compass history stay in this browser. Generated result links contain the data needed to reopen a result; Focus Area links include selected answer IDs. Separate first-party aggregate counters may receive derived buckets when coarse measurement is enabled. The site does not collect research responses or contact details.

Any future research study would require separate information about its purpose, consent, retention, deletion, and access rules.

Read the privacy and data-use note.